

# Curriculum Learning – Group K

This is the homework assignment for Curriculum Learning for group K.

We will show the reward function used to compute the factorized reward. The different rewards used to compute the final one are collision avoidance, progress toward the goal, smooth steering, time to target and goal reaching.

We uploaded the turtlebot3\_world.py file, in which the goal sampling curriculum is also implemented.

## Reward Function

```
def _compute_reward(self, observations, done):
    # --- 1. State and Distance ---
    x, y, _ = self._get_robot_pose()
    dist = self._distance_to_goal(x, y)

    if self.prev_dist is None:
        self.prev_dist = dist

    # --- 2. Compute factorized components ---

    # A. Progress Reward: based on distance difference (Reward Shaping)
    # If diff > 0 the robot got closer, if < 0 it got farther
    diff = self.prev_dist - dist
    r_progress = self.w_progress * diff

    # B. Safety / Collision Avoidance
    # Compute min valid laser scan
    valid_scan = [o for o in observations if not numpy.isinf(o) and not
    numpy.isnan(o)]
    min_scan = min(valid_scan) if len(valid_scan) > 0 else 10.0

    r_collision_avoid = 0.0
    threshold_safe = 0.3 # meters
    if min_scan < threshold_safe:
        # Linear penalty: the closer it is, the more negative the reward
        r_collision_avoid = -self.w_collision * (threshold_safe - min_scan)
    else:
        r_collision_avoid = 0.1 * self.w_collision # Small reward for being in
a safe zone

    # C. Smooth Steering
    # Penalize if last action was not FORWARDS
    r_smooth = 0.0
```

```

if self.last_action != "FORWARDS":
    r_smooth = -self.w_smooth

# D. Terminal Rewards
r_terminal = 0.0
if done:
    if self.reached_goal:
        r_terminal = self.terminal_goal
    elif self.steps >= self.max_steps_per_episode:
        r_terminal = self.terminal_timeout
    else:
        # If done is True but not goal or timeout, it's a collision
        r_terminal = self.terminal_crash

# E. Time Rewards
r_cumm_time = self.r_time * self.steps

# --- 4. Final computation and update ---
reward = r_progress + r_cumm_time + r_smooth + r_collision_avoid +
r_terminal

# Update cumulative reward components
self.cum_r_progress += r_progress
self.cum_r_time += self.r_time
self.cum_r_smooth += r_smooth
self.cum_r_collision_avoid += r_collision_avoid
self.cum_r_terminal += r_terminal
# Update state variables for next step
self.prev_dist = dist
self.prev_action = self.last_action

# --- 5. Logging structure ---
self.last_reward_components = {
    "r_progress": float(r_progress),
    "r_time": float(self.r_time),
    "r_smooth": float(r_smooth),
    "r_collision_avoid": float(r_collision_avoid),
    "r_terminal": float(r_terminal),
    "dist": float(dist),
    "min_scan": float(min_scan),
    "goal_x": float(self.goal_xy[0]),
    "goal_y": float(self.goal_xy[1]),
    "r_total": float(reward),
    "cum_r_progress": self.cum_r_progress,
    "cum_r_time": self.cum_r_time,
    "cum_r_smooth": self.cum_r_smooth,
    "cum_r_collision_avoid": self.cum_r_collision_avoid,
    "cum_r_terminal": self.cum_r_terminal,
    "cumulated_reward": self.cumulated_reward,
}

```

```

        self.cumulated_reward += reward
        self.steps += 1

        print("Step Reward: {:.2f} | Components: {}".format(
            reward, self.last_reward_components))

    return reward

```

## Tests

### Run 1

Deep learning model

```


def __init__(self, inputs, outputs):
    super(DQN, self).__init__()
    self.fc1 = nn.Linear(inputs, 64)
    self.fc2 = nn.Linear(64, 128)
    self.fc3 = nn.Linear(128, 64)
    self.head = nn.Linear(64, outputs)

```

Rewards weights and values

Reward	Value
w_progress	20.0
w_collision	2.0
w_smooth	0.15
terminal_goal	100.0
terminal_crash	-25.0
terminal_timeout	-10.0
r_time	-0.05

## Performance metrics



## Run 2


Deep learning model

```
def __init__(self, inputs, outputs):
    super(DQN, self).__init__()
    self.fc1 = nn.Linear(inputs, 64)
    self.fc2 = nn.Linear(64, 128)
    self.fc3 = nn.Linear(128, 64)
    self.head = nn.Linear(64, outputs)
```

Rewards weights and values

Reward	Value
w_progress	20.0
w_collision	10.0
w_smooth	0.1
terminal_goal	150.0
terminal_crash	-100.0
terminal_timeout	-10.0
r_time	-0.05

## Performance metrics



## Run 3




Deep learning model

```
def __init__(self, inputs, outputs):
    super(DQN, self).__init__()
    self.fc1 = nn.Linear(inputs, 64)
    self.fc2 = nn.Linear(64, 128)
    self.fc3 = nn.Linear(128, 256)
    self.fc4 = nn.Linear(256, 64)
    self.head = nn.Linear(64, outputs)
```

Rewards weights and values

Reward	Value
w_progress	20.0
w_collision	5.0
w_smooth	0.1
terminal_goal	100.0
terminal_crash	-50.0
terminal_timeout	-10.0
r_time	-0.05

## Performance metrics



## Run 4

### Deep learning model

```
class ResidualBlock(nn.Module):
    def __init__(self, dim: int):
        super().__init__()
        self.ln1 = nn.LayerNorm(dim)
        self.fc1 = nn.Linear(dim, dim)
        self.ln2 = nn.LayerNorm(dim)
        self.fc2 = nn.Linear(dim, dim)

    def forward(self, x):
        h = self.fc1(F.relu(self.ln1(x)))
        h = self.fc2(F.relu(self.ln2(h)))
        return x + h

class DQN(nn.Module):
    def __init__(self, inputs, outputs, dim: int = 128, n_blocks: int = 4):
        super(DQN, self).__init__()
        self.stem = nn.Sequential(
            nn.Linear(inputs, dim),
            nn.ReLU(),
        )
        self.blocks = nn.Sequential(*[ResidualBlock(dim) for _ in range(n_blocks)])
        self.head = nn.Linear(dim, outputs)

    # Called with either one element to determine next action, or a batch
    # during optimization. Returns tensor([[left0exp,right0exp]...]).
    def forward(self, x):
        x = x.to(device)
        x = self.stem(x)
        x = self.blocks(x)
        return self.head(x)
```

### Rewards weights and values

Reward	Value
w_progress	20.0
w_collision	5.0
w_smooth	0.1
terminal_goal	100.0
terminal_crash	-50.0
terminal_timeout	-10.0
r_time	-0.05

## Performance metrics

